

THE PROMPT GUIDEBOOK_

How to talk to AI and get exactly what you ask for. The CLEAR framework, 30+ copy-paste recipes, and the ten mistakes almost everyone makes.

FREE · MY GIFT TO YOU

JOSEPH CONLAN

CloneCentre · clonecentre.ai
LIVERPOOL · 2026

> How to Talk to AI and Get Exactly What You Ask For

By Joseph Conlan, Founder of Clone Centre, Liverpool, UK

Inside This Guide

A Letter from Joseph

What a Prompt Really Is

The Anatomy of a World Class Prompt

The Clone Centre CLEAR Framework

The Prompt Recipe Library

Advanced Techniques

The Ten Most Common Prompting Mistakes

Building Your Own Prompt Library

Glossary in Plain English

About the Author and Clone Centre

A Letter from Joseph

Hello, and thank you for picking up this guide.

My name is Joseph Conlan. I run Clone Centre, an AI consultancy based in Liverpool, and I spend my days helping businesses work out what AI can actually do for them. Not the headlines version. The practical, Monday morning version, where a real person sits down with a real problem and wants a useful answer.

Over the past few years I have sat in rooms with billion dollar VC funds while they decided where to place their bets on AI. I have watched founders pitch, engineers build, and boards argue about strategy. And I can tell you the thing that surprised me most, it was never the technology that separated the winners from the rest. The

models are available to everyone. What separated them was something much simpler and much older. It was the ability to say precisely what they wanted.

That is all prompting is. Telling a very capable system exactly what you want, in a way it cannot misunderstand.

Here is what I see every week. Two people sit down with the same AI tool. One types "write me a marketing email" and gets something bland they delete within seconds. The other spends ninety seconds describing who they are, who the email is for, what it must achieve and what it must avoid, and gets a draft they can send after five minutes of editing. Same tool. Same price. Completely different result. The difference is not talent, and it is not luck. It is a learnable skill, and this book teaches it.

I wrote this guide because prompting is the one AI skill that matters for almost everyone. You do not need to code. You do not need to understand the mathematics. You need to know how to brief the machine the way you would brief a brilliant new colleague: clearly, completely, and with confidence that it wants to help.

A few honest promises before we start. I will not drown you in jargon. Where a technical term is genuinely useful, I will explain it in plain English, and there is a glossary at the back. I will not pretend AI is magic or that it never gets things wrong, because it does, and Chapter 2 explains why. And I will not pad these pages with hype. Everything in here is something I use myself or teach to clients.

Here is how to get the most from the book. Chapters 1 to 4 give you the foundations and the framework. Chapter 5 is a recipe library you can copy, paste and adapt from today; keep it within reach of your keyboard. Chapters 6 to 8 take you from competent to genuinely skilled. The glossary and the notes at the back are there whenever you need them.

One last thing. Do not just read this book. Use it. Open your AI tool of choice and try every technique as you meet it. Prompting is a practical skill, like cooking or driving. Ten minutes of doing beats an hour of reading about doing.

Right. Let us begin.

Joseph Conlan Liverpool, July 2026

What a Prompt Really Is

Before you can write great prompts, it helps to know what happens to your words after you press enter. This chapter explains that in plain English. No mathematics, I promise. But the few ideas here will make every other chapter click into place.

> The thirty second version

A large language model, the technology behind tools like ChatGPT, Claude and Gemini, is a prediction engine. It has been trained on an enormous collection of text: books, articles, websites, conversations. From that training it learned the patterns of language. Which words tend to follow which other words, in which situations, in which styles.

When you send it a prompt, it does not look up an answer in a database. It predicts, one piece at a time, what a helpful, plausible response to your words would look like. That is the whole trick. Everything else in this book follows from it.

> It reads chunks, not words

The model does not actually see letters or whole words. It sees tokens: small chunks of text. A token might be a whole word like "cat", part of a word like "ing", or a piece of punctuation. Common words are usually one token. Longer or rarer words get split into several.

Why should you care? Two reasons. First, it explains why AI sometimes fumbles things that seem trivial to us, like counting the letters in a word, because it never saw the letters, only the chunks. Second, it explains why long, rambling prompts cost you. Every word you write is converted into tokens, and the model has to process all of them. Clear and compact beats long and waffly, every time.

> The context window: its entire attention span

A model can only consider a limited amount of text at once. That limit is called the context window, and it includes your prompt, the conversation so far, and the answer it is writing. Modern models have large windows, but they are still finite.

Think of it like a desk. Everything on the desk, the model can see and use. Anything not on the desk does not exist as far as it is concerned. It does not remember your website, your company, or the conversation you had with it last Tuesday. If it is not in the window, it is not in its world.

JOSEPH SAYS

Most "the AI was useless today" stories I hear are really context window stories. The person assumed the tool knew something it had never been told. Once you accept that the model only knows what is in front of it, your prompts improve overnight.

> It predicts; it does not understand

Here is the most important idea in the chapter. The model does not understand your request the way a person does. It has no beliefs, no intentions and no common sense

in the human sense. It is extraordinarily good at producing text that matches the pattern your prompt sets up.

This explains its greatest strength and its greatest weakness.

The strength: because it has seen so much text, it can produce a convincing legal letter, a warm bedtime story, a Python script or a marketing plan, simply because your prompt pointed it at that region of its training. You are not commanding a machine. You are steering a pattern.

The weakness: because it predicts plausible text rather than verified truth, it can state false things with total confidence. This is called hallucination, and it is not a bug that will be fully fixed. It is a natural consequence of how the thing works. If the prompt sets up a pattern where an answer should exist, the model will often produce one whether or not the facts are behind it.

So rule one of working with AI: for anything that must be true, check. Names, dates, figures, legal points, medical points. Treat the model as a brilliant assistant who occasionally bluffs, and you will get along famously.

> Why your exact words matter so much

Because the model is matching patterns, small changes in your prompt cause large changes in the output. Each word you write tilts the probabilities of what comes next.

Tell it "you are a cautious financial adviser" and the entire response shifts: more hedging, more structure, more formality. Tell it "explain this to a tired parent at 10pm" and the same question gets a completely different answer. Neither instruction changed the model. They changed the pattern it was completing.

This is why prompting is a skill rather than a lottery. The levers are real, they are learnable, and they are exactly what the rest of this book is about:

- **Vocabulary** sets the register. Formal words invite formal answers.
- **Structure** invites structure. Numbered requests get numbered answers.
- **Roles** change the knowledge and style the model draws on.
- **Limits** tell it where the edges are, so it stops filling gaps with guesses.

> The polite stranger test

Here is a mental habit I teach every client. Before you send a prompt, read it back and ask: if a capable stranger walked into my office and I handed them only these words, could they do the job? Would they know who the work is for, what good looks like, how long it should be, what to avoid?

If the answer is no, the model cannot do it either. Not because it is not clever enough, but because you have not told it enough. The fix is never "the AI is rubbish". The fix is a better brief.

That brief is what the next chapter is about.

The Anatomy of a World Class Prompt

Every strong prompt I have ever written, for clients or for myself, is built from the same seven parts. You will not need all seven every time. A quick question needs a sentence. But for anything that matters, a proposal, a report, a campaign, a plan, running through these seven is the difference between a throwaway answer and a working document.

Here they are, in the order I think about them:

1. **Role:** who the AI should be.
2. **Task:** what you actually want done.
3. **Context:** the background it cannot guess.
4. **Format:** the shape of the answer.
5. **Tone:** how it should sound.
6. **Constraints:** what it must not do.
7. **Examples:** a model of what good looks like.

Let us take them one at a time. For each one I have shown a weak version and a strong version side by side, so you can see the gap for yourself.

> 1. Role

The role tells the model which part of its training to draw on. "You are a solicitor" and "you are a stand up comedian" produce different vocabularies, different structures and different instincts, because the model has seen how both of those people write.

A good role is specific. Not "you are an expert" but "you are a chartered accountant who specialises in UK small business tax". The more precise the professional identity, the more precise the instincts the model borrows.

✗ WEAK	✓ STRONG
You are an expert.	You are a commercial solicitor in England who reviews supplier contracts for small businesses.
Act like a marketer.	You are a direct response copywriter with twenty years of experience writing emails for independent retailers.
Be a teacher.	You are a patient maths tutor who specialises in helping adults who struggled with maths at school.

One caution. The role is a lens, not a licence. "You are a doctor" will get you text that sounds medical, but it does not make the output medically reliable. The role shapes style and emphasis, not truth.

> 2. Task

The task is the verb. It is the single sentence the whole prompt hangs on, and it is where most weak prompts fail. "Help me with my website" is not a task. "Rewrite my homepage headline so a first time visitor understands what we sell within five seconds" is a task.

Strong tasks have three properties. They use a clear verb (write, summarise, compare, list, draft, critique, rewrite, plan). They name the deliverable (a 200 word email, a table of five options, a one page plan). And they state what success looks like.

✗ WEAK	✓ STRONG
Do something with this report.	Summarise this report in five bullet points, each under twenty words, highlighting anything that needs a decision this week.
Write about our new product.	Write a 150 word product announcement for our new boiler service plan, aimed at existing customers, leading with the price freeze.
Give me ideas.	List ten low cost marketing ideas for a Liverpool coffee shop, ranked from easiest to hardest to implement.

If you can only improve one thing in your prompting, improve your tasks. A precise task rescues a mediocre prompt faster than any other fix.

> 3. Context

Context is everything the model needs to know and cannot possibly guess. Remember the polite stranger test from Chapter 2. Context is how you pass it.

What belongs in context: who you are, who the work is for, what has already happened, what the reader knows already, and any facts or figures the output must use. What does not belong: your life story. Context is not length, it is relevance. Three sharp sentences beat three rambling paragraphs.

✗ WEAK	✓ STRONG
Write a reply to this customer.	Write a reply to this customer. They ordered on 3 May, the parcel arrived damaged on 10 May, and our policy is to replace damaged goods without requiring a return. They have bought from us four times before.
Plan my week.	Plan my week. I work 9 to 5 Monday to Friday, I have fixed meetings on Tuesday and Thursday mornings, my most important project is the Henderson proposal due Friday, and I do my best deep work before 11am.
Improve this email.	Improve this email. It goes to a landlord who has not responded to two previous messages about a broken boiler. I want a reply without damaging the relationship, because my tenancy renews in August.

A useful test for context: imagine the output is wrong. Ask yourself, what fact did the model not have that would have prevented that mistake? That fact belongs in your next prompt.

> 4. Format

Format is the shape of the answer, and it is the cheapest win in prompting. Left alone, the model guesses a shape, and its guess is usually "a few paragraphs of prose". If you wanted a table, a checklist or a script, you will be reformatting by hand.

Tell it the shape explicitly: how many sections, what order, how long, and what structure. You can even show it the skeleton and ask it to fill it in.

✗ WEAK	✓ STRONG
Compare these two suppliers.	Compare these two suppliers in a table with columns for price, delivery time, warranty and payment terms, then give me a two sentence verdict underneath.
Give me a training plan.	Give me a four week training plan. Use one section per week. In each section give me three sessions, and for each session list the exercises, sets and rest times as bullets.
Summarise this article.	Summarise this article in exactly three parts: a one sentence overview, three bullet points of key facts, and one sentence on why it matters for a small business owner.

Format instructions also make outputs consistent. If you ask for the same format every week, your weekly report summary comes back in the same shape every week, which makes your own work faster.

> 5. Tone

Tone is how it sounds, and it matters more than people expect, because the default AI tone is recognisably bland: cheerful, padded and vaguely corporate. Your readers can smell it.

The fix is to describe the voice you want with reference points the model knows. "Warm but professional" is a start. "Like a trusted local accountant explaining something to a long standing client" is far better, because it points at a real pattern of writing.

x WEAK	✓ STRONG
Make it friendly.	Write it in the voice of a plain speaking Yorkshire shop owner: warm, direct, no jargon, short sentences.
Be professional.	Formal but human, the tone of a well written letter from a good solicitor: courteous, precise, no legalese, no exclamation marks.
Make it punchy.	Short sentences. Active verbs. Confident but not arrogant. Think of a founder writing to their own customers, not a corporation issuing a statement.

A trick I use constantly: name what to avoid as well as what to aim for. "No corporate jargon, no buzzwords, no exclamation marks" removes three of the biggest tells of generic AI writing in one line.

> 6. Constraints

Constraints are the fences. They tell the model where it must not go, and they are what stop a good answer turning into an unusable one. Without constraints, the model pads, wanders and fills gaps with confident invention.

The most useful constraints cover length, content and behaviour. Length: "under 200 words". Content: "use only the figures in the table below; do not invent statistics". Behaviour: "if you are unsure of something, say so rather than guessing".

✗ WEAK	✓ STRONG
Write a job advert.	Write a job advert of no more than 250 words. Do not use the words "rockstar", "ninja" or "fast paced". State the salary range I give you exactly as written.
Explain inheritance tax.	Explain inheritance tax in plain English for a complete beginner. Do not give specific financial advice, do not quote thresholds you are unsure of, and flag anything the reader must verify with a professional.
Draft a social post about our results.	Draft a social post about these results. Use only the numbers in my message. If a number is missing, leave it out rather than estimating. Keep it under 80 words.

That last example hides the single most valuable constraint in business prompting: use only what I gave you. It will not eliminate every error, but it dramatically reduces confident invention.

> 7. Examples

Examples are the most powerful and the most underused part of the anatomy. One example of what you want is worth three paragraphs of description, because the model is a pattern matcher. Show it the pattern and it will follow it.

An example can be a sample of your own writing to copy the voice of, a previous document in the right format, or a made up model answer showing the structure you want. Paste it in, label it clearly, and say "match this".

✗ WEAK	✓ STRONG
Write product descriptions in our style.	Here are two of our existing product descriptions: [paste]. Write the new one so a customer could not tell a different person wrote it. Match the sentence length, the humour and the structure.
Format the addresses properly.	Format each address like this example: "Mrs J. Smith, 14 Lark Lane, Liverpool L17 8UU". Here are the addresses: [paste].
Give feedback on my essay.	Here is feedback you gave me last time that was exactly right: [paste]. Give feedback on this new essay in the same format, with the same balance of praise and specific criticism.

> Putting the anatomy together

Here is a prompt that uses all seven parts. Notice how ordinary the language is.

There is no secret syntax. It is simply a complete brief.

THE CLONE CENTRE PROMPT GUIDEBOOK • CLONECENTRE.AI

You are a marketing consultant who specialises in UK independent coffee shops. [ROLE]

Write a one page promotional plan for the next month. [TASK]

I run a small coffee shop on Lark Lane in Liverpool. We are busy on weekend mornings but quiet on weekday afternoons. Our margin on food is better than on drinks. We have no marketing budget beyond about £100 and two hours a week of staff time. [CONTEXT]

Structure it as: the single main goal, three actions, what each action costs in money and time, and how we will know it worked. [FORMAT]

Plain speaking and practical, like a consultant a friend recommended. No buzzwords. [TONE]

Do not suggest anything involving paid advertising or discounts over 10 percent.
Do not invent local statistics. [CONSTRAINTS]

Here is a plan I liked from a previous month, match its structure and level of detail:
[pasted example]. [EXAMPLES]

TRY THIS NOW

Take a task you actually need done this week and write it out as a seven part brief like the one above. It will feel slow the first time. By the third time it takes two minutes, and the quality gap will astonish you.

You do not need to label the parts in real prompts; I did it here so you can see the skeleton. In practice it reads as one natural paragraph or two. But the skeleton is always there in my best prompts, and now it can be in yours.

The next chapter turns this anatomy into a framework you can run through in under a minute: CLEAR.

The Clone Centre CLEAR Framework

The seven part anatomy from Chapter 3 is complete, but in the middle of a busy day you will not recite seven things. So at Clone Centre we teach a shorter checklist that captures the parts people actually forget. It is called CLEAR, and it takes under a minute to run through.

LETTER	STANDS FOR	THE QUESTION TO ASK YOURSELF
C	Context	Have I told it what it cannot guess?
L	Limits	Have I set the fences: length, scope, what to avoid?
E	Examples	Have I shown it what good looks like?
A	Audience	Have I said who this is for?
R	Review	Have I told it how I will judge the answer?

CLEAR is not a replacement for the anatomy. Role, Task and Tone still matter. Think of it this way: the task is the engine, and CLEAR is the pre flight check that stops you forgetting the fuel.

Let me walk you through each letter, then we will build a real prompt together, one letter at a time.

> C: Context

Ask: what does the model need to know that it cannot possibly guess? Your situation, the background, the facts the answer must use. Chapter 3 covered this in depth, so here I will just add the habit: before sending, scan your prompt and delete any sentence starting "you probably know" or "as you can see". It does not know and it cannot see. Put the fact in.

> L: Limits

Ask: where are the edges? Length limits ("under 150 words"), content limits ("only use the figures I gave you"), style limits ("no jargon"), and scope limits ("answer the question; do not draft the email yet"). Limits are what turn an impressive answer into a usable one.

> E: Examples

Ask: can I show rather than tell? A sample paragraph of your own writing, a document in the right format, a model answer. If you have nothing to paste, you can describe a pattern ("use short sentences like a text from a friend"), but a real example beats a description every time.

> A: Audience

Ask: who is this actually for? The same content lands completely differently depending on the reader. "Explain our price rise" for a long standing trade customer

is a different document from the same explanation for a new retail customer. Always name the audience and, where it helps, name what they care about.

> R: Review

Ask: have I told it how I will judge success? This is the letter people skip, and it has two parts. First, tell the model the standard before it answers ("I will judge this on whether a busy manager can act on it in under a minute"). Second, review the answer yourself against that standard and go back with specific corrections. Prompting is a loop, not a single shot.

WATCH OUT

The R is where average users and strong users separate. Average users accept the first answer or abandon the tool. Strong users treat the first answer as a draft and reply with precise corrections: "shorter", "less formal", "you invented a figure in paragraph two, remove it". Two rounds of review beat ten minutes of prompt crafting.

> A worked example, built step by step

Let me show you CLEAR in action. Imagine I run a landscaping business. A potential customer, Sarah, asked for a quote two weeks ago. I sent it, and I have heard nothing since. I want a follow up email.

Here is how most people prompt it:

```
Write a follow up email about the landscaping quote.
```

It will produce something. It will be generic, slightly desperate, and it will mention nothing Sarah actually cares about. Now watch what happens as we apply each letter.

Step 1: add Context.

```
Write a follow up email about a landscaping quote.
```

```
Context: I run a small landscaping business in Liverpool. Two weeks ago I visited Sarah's home in Allerton and quoted £4,800 to redesign her back garden: new lawn, raised beds and a porcelain patio. She was enthusiastic on the day and said she wanted it done before her daughter's birthday party in August. I have heard nothing since sending the written quote.
```

Already better. The model now knows the price, the scope, the location and the deadline that actually matters to Sarah.

Step 2: add Limits.

Write a follow up email about a landscaping quote.

Context: I run a small landscaping business in Liverpool. Two weeks ago I visited Sarah's home in Allerton and quoted £4,800 to redesign her back garden: new lawn, raised beds and a porcelain patio. She was enthusiastic on the day and said she wanted it done before her daughter's birthday party in August. I have heard nothing since sending the written quote.

Limits: under 120 words. No pressure and no discount offer. Do not use the phrase "just checking in" or "circling back". One clear question at the end.

Now it cannot be pushy, cannot waffle, and cannot reach for the two most tired follow up phrases in business.

Step 3: add an Example.

Write a follow up email about a landscaping quote.

Context: I run a small landscaping business in Liverpool. Two weeks ago I visited Sarah's home in Allerton and quoted £4,800 to redesign her back garden: new lawn, raised beds and a porcelain patio. She was enthusiastic on the day and said she wanted it done before her daughter's birthday party in August. I have heard nothing since sending the written quote.

Limits: under 120 words. No pressure and no discount offer. Do not use the phrase "just checking in" or "circling back". One clear question at the end.

Here is a follow up I sent last year that won the job. Match its warmth and its structure: short opening reference to the visit, one useful piece of information, one simple question.

[pasted example email]

The model now has a proven pattern to follow rather than inventing one.

Step 4: add the Audience.

Write a follow up email about a landscaping quote.

Context: I run a small landscaping business in Liverpool. Two weeks ago I visited Sarah's home in Allerton and quoted £4,800 to redesign her back garden: new lawn, raised beds and a porcelain patio. She was enthusiastic on the day and said she wanted it done before her daughter's birthday party in August. I have heard nothing since sending the written quote.

Limits: under 120 words. No pressure and no discount offer. Do not use the phrase "just checking in" or "circling back". One clear question at the end.

Here is a follow up I sent last year that won the job. Match its warmth and its structure: short opening reference to the visit, one useful piece of information, one simple question.

[pasted example email]

Audience: Sarah is a busy working parent, not a gardener. She cares about the party deadline and about not being chased. She will read this on her phone between meetings.

Now the email will be short enough for a phone screen, and it will lead with the thing Sarah cares about, the deadline, not the thing I care about, the quote.

Step 5: add the Review standard.

Write a follow up email about a landscaping quote.

Context: I run a small landscaping business in Liverpool. Two weeks ago I visited Sarah's home in Allerton and quoted £4,800 to redesign her back garden: new lawn, raised beds and a porcelain patio. She was enthusiastic on the day and said she wanted it done before her daughter's birthday party in August. I have heard nothing since sending the written quote.

Limits: under 120 words. No pressure and no discount offer. Do not use the phrase "just checking in" or "circling back". One clear question at the end.

Here is a follow up I sent last year that won the job. Match its warmth and its structure: short opening reference to the visit, one useful piece of information, one simple question.

[pasted example email]

Audience: Sarah is a busy working parent, not a gardener. She cares about the party deadline and about not being chased. She will read this on her phone between meetings.

Before you write, note that I will judge this email on one thing: does it give Sarah a genuinely useful reason to reply, rather than just reminding her I exist? If the most helpful hook is the August party deadline and my booking calendar filling up, use it. After the email, give me one sentence on the approach you chose.

Read that final version again. Total time to write: about three minutes. The email it produces will reference the visit, mention the party deadline, note that the build calendar is filling, and ask one simple question, in my voice, at the right length, with no cringe. That is the CLEAR difference.

> The CLEAR pre flight check

Before you send any prompt that matters, ask five questions:

1. **Context:** does it know what it cannot guess?
2. **Limits:** are the fences up?
3. **Examples:** have I shown it good?
4. **Audience:** have I named the reader?
5. **Review:** have I set the standard, and will I send corrections?

Five questions, under a minute, and your hit rate will roughly transform. This is the framework we teach in every Clone Centre engagement, and it has never let me

down.

The Prompt Recipe Library

This is the chapter you will come back to. Thirty two recipes I use myself or hand to clients, organised by the job they do. Copy them, paste them, and replace anything in [square brackets] with your own details.

Three tips before you start. First, always fill in the brackets; a recipe with the brackets left empty is a generic prompt, and you know now what those produce. Second, treat every answer as a first draft and reply with corrections. Third, when you adapt a recipe and it works brilliantly, save your version. Chapter 8 shows you how to build that into a personal library.

> Writing and Marketing

1. The Blog Post Brief

You are a content writer for small businesses. Create a detailed brief for a blog post titled "[title]". My business is [what you do] and the reader is [who they are]. Give me: the working title, a one sentence promise to the reader, a suggested structure of 5 to 7 sections with two bullet points under each, and the single action the reader should take at the end. Keep the tone [tone, e.g. plain and friendly]. Do not write the post yet.

Why it works: separating the brief from the writing gives you a chance to fix the structure before a single paragraph is wasted.

2. Rewrite It In My Voice

Here is a sample of my writing so you can learn my voice: [paste 2 or 3 paragraphs you wrote]. Now rewrite the text below so it sounds like I wrote it. Keep every fact exactly as it is. Match my sentence length, my level of formality and my humour. Do not add new claims.

Text to rewrite: [paste]

Why it works: the model copies patterns, and a real sample of your writing is a far stronger instruction than any description of your style.

3. The Honest LinkedIn Post

Write a LinkedIn post about [topic or lesson]. I am a [your role] and my audience is [who follows you]. Start with a short first line that earns the second line. Tell one specific story or make one specific point, not three. Under 150 words. Plain English, no hashtags, no buzzwords, no "I am humbled to announce". End with a genuine question.

Why it works: banning the platform clichés and forcing one point is what separates posts people read from posts people scroll past.

4. The Newsletter That Gets Read

Draft a monthly email newsletter for [your business]. It goes to [describe your subscribers]. This month's material: [paste your news, notes or bullet points]. Structure: a two sentence personal opening from me, three short sections with a headline each, and one clear call to action: [the action]. Under 300 words total. Warm and direct, like writing to one person, not a list. Flag anything where you need me to fill in a fact.

Why it works: giving it raw material plus a fixed structure means you get a newsletter instead of a generic template with placeholders.

5. Product Description With a Conscience

Write a product description for [product]. It is for [type of customer] buying on [where they buy, e.g. our website]. Facts you must use and must not change: [paste specs, price, materials, sizes]. Structure: one sentence hook about the benefit, three bullet points of features translated into benefits, one sentence on who it is not for. Under 120 words. No hype words like "revolutionary" or "ultimate".

Why it works: the "who it is not for" line builds trust and the banned word list keeps the copy sounding human.

6. Headline Variations With a Point of View

Give me 10 headline options for [the piece: blog, email subject line, advert]. The piece is about [one sentence summary] and the reader is [who]. 5 should promise a clear benefit, 3 should create curiosity without clickbait, and 2 should be plain and direct. Under 12 words each. Then tell me which two you would test first and why.

Why it works: forcing variety across styles plus a reasoned shortlist beats asking for "10 catchy headlines" and getting ten versions of the same one.

7. One Idea, Five Captions

The core idea: [describe your idea, offer or story in 2 or 3 sentences]. Turn it into 5 social captions: one for LinkedIn (professional, 100 words), one for Instagram (conversational, 50 words, line breaks), one for Facebook (warm, community feel), one for X (under 240 characters), and one for Google Business Profile (factual, local, mentions [your town]). Keep the facts identical across all five.

Why it works: one source of facts with five platform specific rewrites keeps your message consistent while respecting how each platform reads.

> Business and Strategy

8. The One Page Business Plan

Help me stress test a business idea. The idea: [describe it in 2 or 3 sentences]. I am based in [location] and can invest [rough time and money]. Ask me the 8 most important questions you would need answered to judge this idea, one at a time. After I answer, produce a one page plan: the customer, the problem, the offer, how I reach the first 10 customers, the main cost, the biggest risk, and the first three actions for this month.

Why it works: making the model interview you first forces the gaps in your thinking into the open before it writes anything.

9. Competitor Teardown

You are a strategy consultant. Here is what I know about my competitor: [paste their website copy, pricing page or your notes]. My business is [what you do and who for]. Analyse: what they do well, what they promise that I also offer, what they miss that I could own, and how their pricing positions them. Present it as a table plus three specific moves I could make in the next 90 days. Only use the information I gave you; say "unknown" where you would otherwise be guessing.

Why it works: the "only use what I gave you" limit turns the model from a rumour machine into a disciplined analyst.

THE CLONE CENTRE PROMPT GUIDEBOOK · CLONECENTRE.AI

10. An Honest SWOT

Run a SWOT analysis for [your business or project]. What we do: [describe]. Our situation: [describe market, size, stage]. Be brutally honest rather than encouraging. Give me 4 points per quadrant, each one specific enough to act on, not generic filler like "competition exists". Then name the single weakness most likely to sink us if ignored, and the single strength we underuse most.

Why it works: asking for brutal honesty and forcing specificity stops the model writing the flattering, vague SWOT it defaults to.

11. Meetings That Earn Their Keep

Create an agenda for a [length] meeting about [purpose] with [who attends]. The decision this meeting must produce is [the decision]. Structure: 5 minutes of context, timed discussion items each with an owner and the outcome needed, and 5 minutes to confirm actions. After the agenda, draft the follow up email template I will send afterwards with blanks for the decisions and actions.

Why it works: defining the required decision up front makes the model build a meeting around an outcome instead of a chat around a topic.

12. The Proposal Skeleton

Draft the structure and key content for a proposal I am sending to [client type] for [the work]. What I know about them: [their situation and what they asked for]. My price is [price or range] and the work takes [timeframe]. Structure it as: their problem restated better than they said it, my approach in 3 phases, what they get at each phase, investment, why me in two sentences, and next step. Write the full text for the problem section only; bullet the rest for me to complete.

Why it works: restating the client's problem well is what wins proposals, so the recipe spends the model's effort there and leaves the boilerplate to you.

13. Pricing Sense Check

I want to pressure test my pricing. I sell [product or service] to [customer type] at [price]. It costs me [cost] to deliver and takes [time]. Alternatives my customer has: [list them]. Do three things: identify what my price implies about my positioning, list the 5 strongest objections a buyer might raise at this price, and suggest 3 ways to increase the perceived value without cutting the price. Be direct; do not just reassure me.

Why it works: separating positioning, objections and value building gives you three usable outputs instead of a vague "your price seems reasonable".

14. The Difficult Conversation Script

I need to have a difficult conversation with [who: employee, supplier, client] about [the issue, in 2 or 3 sentences]. My goal is [what a good outcome looks like] and I want to preserve the relationship because [reason]. Write a loose script: how to open in one sentence, the facts I should state without blame, the question I should ask to hear their side, and two phrases to avoid because they would escalate it. Then list the 3 ways they might react and a calm response to each.

Why it works: rehearsing the other person's likely reactions is what actually lowers the temperature, and the model is a tireless rehearsal partner.

> Finance and Admin

15. The Invoice Chase That Keeps the Client

Write an email chasing an overdue invoice. Facts: invoice [number] for [amount] was due on [date] and is now [X days] overdue. The client is [describe: new, long standing, usually reliable]. This is my [first / second / third] reminder. Firm but courteous, assume it is an oversight, not bad faith. Include the invoice details, ask for a payment date, and mention the original work is complete and they were happy with it. Under 100 words. No threats at this stage.

Why it works: naming the relationship stage and banning threats gets you paid without burning the client.

16. The Plain English Cashflow Picture

Here are my numbers for the next three months: [paste income expected, regular outgoings, one off costs, current balance]. Explain my cashflow position in plain English: where the tight spots are, which weeks to watch, and the 3 actions that would most improve the picture. Use a simple table to show the months. Do not give financial advice or assume anything beyond the numbers I provided; tell me what to check with my accountant.

Why it works: confining the model to your real numbers and forcing it to flag accountant territory keeps it useful and safe.

17. The Sensible Expense Policy

Draft a one page expenses policy for my business. We are [size and type of business]. It should cover: travel, meals, software subscriptions and client entertainment. For each, state what is allowed, what needs approval first and from whom, and what is never claimable. Plain English, no legal jargon, firm but not suspicious in tone. End with the claim process in 4 steps. Flag anything I should confirm with my accountant before issuing this.

Why it works: the "allowed, approval, never" structure is exactly how good policies read, and the flag keeps you compliant.

18. The Spreadsheet Translator

I am working in [Excel / Google Sheets]. My data looks like this: [describe columns and a couple of example rows]. I want to [what you want to calculate or produce]. Give me the exact formula, explain in one sentence how it works, and tell me where to put it. If there are two ways to do it, give me the simpler one. Assume I know basic spreadsheets but nothing advanced.

Why it works: describing your actual columns and asking for the exact formula turns the model into a patient expert sitting next to you.

19. The Contract Translator

Here is a contract I have been asked to sign: [paste the text]. Summarise it in plain English: what I am agreeing to, what they are agreeing to, the payment terms, how either side can end it and with what notice, and any clause that is unusual or one sided. Present unusual clauses as a numbered list with the clause quoted and the concern in one sentence each. You are not my solicitor; end by telling me what a professional should check before I sign.

Why it works: quoting the exact clause next to each concern lets you verify the summary instead of trusting it, which is the only safe way to use AI on legal text.

20. The Monthly Report in Five Minutes

Here are this month's raw numbers and notes: [paste]. Turn them into a short monthly report for [audience: me, my business partner, my board]. Structure: the headline of the month in one sentence, a small table of the key numbers against last month, what went well, what went badly, and the one decision needed next month. Neutral, factual tone. Do not smooth over bad numbers. Under 250 words.

Why it works: a fixed structure plus an instruction not to smooth over bad news gives you an honest report instead of a cheerful summary.

> Learning and Research

21. Explain It Like I Am New

Explain [topic] to me from scratch. Assume I am intelligent but know nothing about this field. Start with the one sentence version, then the paragraph version, then the full explanation. Use one everyday analogy. Define every technical term the first time you use it. End with 3 questions I could ask you next to go deeper. Do not use jargon to explain jargon.

Why it works: the three zoom levels plus a next step menu turns one answer into a whole learning path you control.

22. The Socratic Tutor

Be my tutor on [topic]. Do not lecture me. Ask me one question at a time to find out what I already understand, then teach to the gap. When I get something wrong, do not just correct me; ask the question that helps me spot the mistake myself. Keep your messages short. Start with your first question now. My goal is to be able to [what you want to be able to do].

Why it works: forcing the model to ask rather than tell converts it from a reference book into a tutor, which is how adults actually learn.

23. The Realistic Study Plan

Build me a study plan. I want to learn [skill or subject] well enough to [specific goal]. I can study [hours] per week, on [which days], starting [date]. My current level: [honest description]. Structure the plan by week, with each week having one clear focus, one practical exercise and a way to check I have actually learned it. Keep it realistic; do not fill every day. Include one milestone test at the end of each month.

Why it works: demanding a check on learning each week stops you collecting information and starts you building skill.

24. Summarise Without the Bleach

Summarise this text: [paste article, report or transcript]. Give me: the main argument in two sentences, the 5 most important points as bullets keeping the author's actual meaning, anything surprising or counterintuitive, and one sentence on what the author might be missing. Do not add outside information. If parts are unclear or contradictory, say so instead of tidying them up.

Why it works: banning outside information and asking for contradictions keeps the summary faithful instead of confident and wrong.

25. Argue Both Sides

The question: [the decision or debate, e.g. "should our shop open on Sundays?"]. My context: [your situation]. First, make the strongest possible case for option A in three points. Then the strongest possible case for option B in three points. Then identify the two facts that would most change the answer if I knew them, and tell me how I could find them. Do not pick a side for me.

Why it works: asking for the strongest version of both cases plus the decisive facts beats asking "what should I do" and getting a coin flip dressed as advice.

26. Quiz Me Properly

Quiz me on [topic or pasted material]. Ask 10 questions, one at a time, starting easier and getting harder. After each of my answers tell me: right or wrong, the ideal answer in two sentences, and one related point worth knowing. Wait for my answer before asking the next question. At the end, tell me my 3 weakest areas and what to revise. Here is my material if I am providing it: [paste, or say "your choice of standard questions"].

Why it works: one question at a time with feedback turns the model into an examiner, and the weakest areas report tells you exactly where to spend your next hour.

> Everyday Life

27. The Meal Plan That Survives Tuesday

Plan my dinners for [number] nights. Household: [who eats, including ages]. We like [list a few cuisines or dishes] and dislike [list]. Constraints: no more than [X minutes] cooking on weeknights, budget around [amount], [any dietary needs], and I want leftovers for lunch twice. Give me: the plan as a simple table, one combined shopping list grouped by shop section, and which nights to cook double.

Why it works: the real constraints, time, budget, leftovers, are what separate a plan you follow from a plan you abandon by Wednesday.

28. The Sensible Trip Planner

Plan [number] days in [place] for [who is going]. Our style: [e.g. walk a lot, museums over shopping, eat where locals eat]. Budget: [rough daily budget]. We are based at [area or hotel type]. Build a day by day outline: one anchor activity per day, a nearby place for lunch and dinner, and one rainy day backup per day. Do not cram more than three things into a day. Flag anything that needs booking ahead.

Why it works: limiting each day to one anchor and forcing booking warnings produces a holiday instead of an endurance event.

29. The Complaint Letter That Gets Results

29. The Complaint Letter That Gets Results | GPT-4 Book Prompts | RECENTRE.AI

Write a complaint letter. The facts: [what you bought, when, from whom, what went wrong, and what you have already tried]. I want [refund, replacement, repair]. Firm, factual and polite; no anger, no sarcasm, no legal threats at this stage. Structure: what I bought, what went wrong, what I have tried, what I want and by when. Under 200 words. End with a clear statement of my next step if unresolved.

Why it works: calm facts plus a stated next step gets escalated, while angry letters get filed.

30. Gift Ideas That Are Not Candles

Suggest 8 gift ideas for [person: relationship, age]. What I know about them: [interests, hobbies, recent life events, things they complain about]. Budget up to [amount]. They already own [anything relevant]. No generic suggestions like candles or socks unless there is a specific twist. For each idea, one sentence on why it fits this person specifically, and roughly where I would find it.

Why it works: feeding it the person's actual quirks and banning generic answers forces ideas that feel considered rather than search results.

31. The Decision Helper

Help me think through a decision. The decision: [describe it]. My options: [list them]. What matters most to me: [list 3 or 4 factors, e.g. cost, time, stress, long term]. Build a simple comparison table of the options against my factors. Then ask me the 3 questions whose answers would most affect the choice. Do not decide for me; after I answer, help me see which option my own answers point to.

Why it works: making the model interrogate your priorities before comparing options surfaces what you actually value, which is where good decisions come from.

32. The Week That Fits Real Life

Plan my week. Fixed commitments: [list with days and times]. Must do tasks: [list with rough durations]. Would like to do: [list]. My energy pattern: [e.g. sharp in mornings, useless after 9pm]. Rules: no more than [X] hours of planned work per day, leave [evenings or one day] free, group small admin into one block. Present it as a day by day schedule, and flag anything that does not fit honestly instead of pretending it will.

Why it works: telling it your energy pattern and forcing it to admit what does not fit produces a schedule you will actually keep.

JOSEPH SAYS

Notice what none of these recipes contain: magic words. There is no secret phrase that unlocks better AI. There is only clarity, context, limits and examples, applied over and over. If a recipe here stops working for you one day, do not hunt for a new spell. Check the brief against CLEAR, fix what is missing, and it will work again.

Advanced Techniques

You now have the anatomy and the CLEAR framework, which cover ninety percent of good prompting. This chapter adds six techniques for the other ten percent: the hard problems, the high stakes documents and the moments when the first answer is not good enough. None of them require anything technical. They are habits of thought, written down.

> 1. Chain of thought: make it show its working

Ask a model "what should I charge for this service?" and it will give you a number with little to stand on. Ask it to reason first, and the answer improves dramatically, because each step of reasoning becomes context for the next, and the conclusion inherits the logic instead of replacing it.

The technique is simple: ask the model to think through the problem step by step before giving its final answer.

Work out a fair day rate for my freelance design work. Think step by step: first list the costs my rate needs to cover, then work out how many billable days I realistically have in a year, then consider what the market position implies, and only then propose a rate. Show your reasoning at each step, then give me the final figure and its assumptions in a short summary.

Use it for anything involving judgement, calculation or trade offs: pricing, planning, comparing options, estimating. Two cautions. First, showing its working does not make the working correct; check the steps. Second, for factual lookups this does nothing, because reasoning cannot fix missing facts.

> 2. Few shot examples: teach the pattern, not the rule

In Chapter 3 we used one example to set a pattern. Few shot prompting is the same idea with several examples, and it is the single most reliable way to get exactly

the output shape you want, especially for repetitive jobs like categorising, formatting or scoring.

```
Classify each customer message as BILLING, TECHNICAL, SALES or OTHER, and rate the urgency LOW, MEDIUM or HIGH. Follow this pattern exactly:
```

```
Message: "I was charged twice this month"
```

```
Category: BILLING, Urgency: HIGH
```

```
Message: "How do I reset my password?"
```

```
Category: TECHNICAL, Urgency: LOW
```

```
Message: "Do you offer discounts for annual plans?"
```

```
Category: SALES, Urgency: MEDIUM
```

```
Now classify these:
```

```
[paste messages]
```

Three examples usually beat one, because the model can see what varies and what stays fixed. Keep the examples realistic and consistent with each other; one contradictory example confuses more than it teaches.

> 3. Iterative refinement: treat the first answer as a draft

This is the least glamorous technique and the one professionals use most. Never expect a finished product from one message. Expect a draft, then edit it by conversation.

The loop looks like this:

1. Send your CLEAR prompt.
2. Read the answer against your standard.
3. Reply with specific corrections: "halve the length", "the second section is too formal, rewrite it like the first", "you invented a statistic in paragraph three, remove it and use only my figures".
4. Repeat until it is right, usually two or three rounds.

Two habits make this work. Be specific in corrections: "make it better" teaches the model nothing, while "cut every sentence over fifteen words" teaches it exactly what you want. And once the output is right, save the whole exchange. That conversation is now a reusable asset (Chapter 8).

> 4. Persona stacking: borrow two brains at once

One role is good. Two complementary roles are better for work that sits at a boundary, because each persona pulls the answer towards a different kind of quality.

Answer this with two personas working together: a cautious compliance officer who spots risk, and a practical operations manager who has to keep the business moving. Draft the policy so it satisfies them both. Where they would disagree, note the tension in brackets rather than hiding it.

Useful pairings I use: copywriter plus sceptical customer, accountant plus salesperson, engineer plus end user. The tension between the personas is where the quality lives. Do not stack more than two or three; past that the voices blur into the generic middle you were trying to escape.

> 5. Negative constraints: the power of "do not"

Models default to certain habits: padding, hedging, buzzwords, invented figures, excessive bullets, over-apology. Negative constraints are how you switch those habits off, and for business writing they often matter more than the positive instructions.

The most valuable negative constraints in my own library:

- "Do not invent statistics, sources or quotes. If a fact is missing, say so."
- "Do not use buzzwords: no 'revolutionary', 'game changing', 'unlock' or 'leverage'."
- "Do not start with 'In today's fast paced world' or any equivalent filler opening."
- "Do not hedge every sentence. One caveat where it matters, not one per paragraph."
- "Do not summarise what I said back to me. Start answering immediately."

A word of caution: negative constraints limit, they do not direct. "Do not be boring" bans one failure but suggests nothing better. Pair every "do not" with a positive instruction about what to do instead.

> 6. Temperature and system prompts: the two dials behind the curtain

If you use AI through an app, you may never see these. If you use a developer tool, an API or a business platform, you will, and it is worth knowing what they do.

Temperature is a dial for predictability. At low temperature, the model picks the most likely next word almost every time: consistent, safe, a bit flat. At high temperature it takes more chances: more surprising, more creative, more likely to wander. Rule of thumb: factual and structured work wants low temperature, brainstorming and creative work can handle more. If you never touch it, the default is a sensible middle.

System prompt is a layer of instruction that sits above your conversation, set by whoever runs the tool. It defines the model's standing personality and rules: "you are the support assistant for X company; be concise; never discuss competitors". Your messages then work within those rules. If you ever build anything with AI, this is where the real control lives, and it is written with exactly the same anatomy you learned in Chapter 3: role, task, context, format, tone, constraints, examples.

TRY THIS NOW

Take one prompt you use regularly and add the five negative constraints from the list above. Compare the output with what you got before. For most business writing this single change removes the recognisably "AI" flavour in one move.

The Ten Most Common Prompting Mistakes

I review a lot of other people's prompts, in client workshops and in my own inbox. The same ten mistakes account for nearly every disappointing result I see. Here they are, with the fix for each.

> Mistake 1: The vague task

"Help me grow my business" gives the model nothing to grip. It will respond with the average of everything it has ever read about business growth, which is to say, filler.

The fix: one verb, one deliverable, one success test. "List five ways a Liverpool dog groomer could fill Tuesday and Wednesday slots, ranked by cost, each with a first step" is a task the model can actually complete.

> Mistake 2: Assuming it knows your world

The model does not know your prices, your customers, your town or what happened in yesterday's meeting. When the output ignores them, the model is not being stupid. You briefed it like a stranger and it behaved like one.

The fix: run the polite stranger test. Put the three facts the answer most depends on directly into the prompt. If a number matters, paste the number.

> Mistake 3: Asking for everything at once

"Research the market, write the business plan, create the marketing strategy and draft the website copy" produces four shallow things instead of one good one. Long multi-part asks dilute attention, exactly as they would with a person.

The fix: sequence your asks. Do the research, review it, then the plan, review it, then the copy. A conversation is free; use the turns.

> Mistake 4: Accepting the first answer

The first answer is the model's average response to your words. It is a draft. People who accept it as final are doing the editing job badly or not at all.

The fix: never send or use a first draft. Reply with at least one round of specific corrections. Two rounds of refinement is where the quality lives.

> Mistake 5: No format, so it picks one

Ask nothing about shape and you get prose paragraphs, because that is the centre of gravity of internet text. Then you spend ten minutes restructuring by hand.

The fix: name the shape every time: a table, five bullets, a 150 word email, a three part structure. Format instructions cost you one sentence and save you reformatting forever.

> Mistake 6: Forgetting the audience

A price rise letter for a loyal trade customer and the same letter for a price comparison shopper are different documents. Write it for "customers" and it lands with nobody.

The fix: always name the reader and what they care about. One sentence: "this is for landlords who have worked with us for years and value reliability over price" changes every word of the output.

> Mistake 7: Trusting facts without checking

The model predicts plausible text. It does not verify. It will state a wrong figure with the same confidence as a right one, because confidence is a style, not a measurement.

The fix: for anything that must be true, names, dates, figures, law, health, tax, check against a primary source. Add the constraint "use only the facts I gave you" wherever you can, and treat the rest as a draft to verify, not a record to trust.

> Mistake 8: Writing an essay of a prompt

The overcorrection. Once people learn context matters, some start writing prompts longer than the output they want, full of repetition and backstory. Buried in noise, the model loses the thread just as a person would.

The fix: every sentence in your prompt should earn its place. If removing a sentence would not change the output, remove it. Dense and relevant beats long and thorough.

> Mistake 9: Describing instead of showing

"Professional but friendly with a hint of wit" is a vibe, and the model's interpretation of it will not be yours. People describe for paragraphs what one example would pin down exactly.

The fix: paste an example. Your own writing, a document you admire, a previous answer that was right. One example outweighs a paragraph of adjectives, every time.

> Mistake 10: Not keeping what worked

The quietest and costliest mistake. Someone crafts a brilliant prompt, gets a brilliant result, and it vanishes into their chat history. Next month they start from scratch and wonder why the magic is gone.

The fix: the moment an output is good enough to use, save the prompt that produced it. That is an asset. The next chapter shows you how to build those assets into a library that compounds.

Building Your Own Prompt Library

Everything in this book compounds if you keep it. A prompt that produced one good document will produce another next month, and a better one the month after, because you will improve it a little each time. That growing collection is your prompt library, and after a year of keeping one, it becomes one of the most useful business assets you own. It is how I work, and it is how every efficient AI user I know works.

Here is the whole system. It is deliberately simple, because a complicated system you abandon is worth less than a simple one you keep.

> Where to keep it

Anywhere you will actually look. A notes app, a spreadsheet, a folder of text files, THE CLONE CENTRE PROMPT GUIDEBOOK : CLONECENTRE.AT a private Git repository if you are that way inclined. The tool does not matter.

What matters is three properties: it is searchable, it is always with you, and adding to it takes under a minute. If saving a prompt is a chore, you will not save them, and the library dies.

> What a library entry looks like

Each saved prompt gets a small record around it. My own template has five fields:

FIELD	WHAT GOES IN IT
Name	A short handle you will search for: "Invoice chase, first reminder"
The prompt	The full text, with [brackets] marking the parts to fill in each time
Version	v1, v2, v3. Bump it every time you improve the wording
Last used and result	Date, and one line on how well the output worked
Notes	Anything to watch: "tone goes stiff if the client is new, soften line two"

That is it. Five fields, a minute to fill in, and the prompt is never lost again.

> Versioning: how prompts get better

Treat your prompts like recipes in a professional kitchen. The first version is a starting point. Every time you use one and the output disappoints in some way, do not just fix the output. Fix the prompt.

The loop is simple:

1. Use the prompt.
2. Notice the correction you had to make ("too formal", "invented a figure", "too long").
3. Add that correction to the prompt itself as a constraint, so you never have to make it again.
4. Save it as the next version.

Six months of this and your top ten prompts are extraordinary, because they carry every lesson you ever learned using them. This is also why downloaded prompt packs disappoint: they carry nobody's lessons. Your library is valuable precisely because it is yours.

> Testing before you trust

Before a prompt graduates to regular use, run it against two or three different inputs. A prompt that works once might be lucky. A prompt that works on three different clients, topics or months is a tool. Keep the test inputs in the notes field; they double as examples when you revisit the prompt months later.

> Sharing with a team

If you work with others, the library is even more valuable shared, because it spreads quality and consistency at the same time. Two rules make sharing work. First, every shared prompt keeps its owner: the person who improves it and answers questions about it. Second, everyone can propose changes, but the owner merges them. A library with ten unedited authors decays into a junk drawer within a quarter.

JOSEPH SAYS

Start tonight with five prompts. The five you actually used this month, cleaned up, bracketed and saved. Do not build a cathedral on day one. A small library you use beats a grand one you curate, and it will grow on its own from here.

Glossary in Plain English

Every technical term in this book, plus the handful you will meet elsewhere, explained the way I would explain them across a table.

TERM	WHAT IT ACTUALLY MEANS
AI (artificial intelligence)	Software that performs tasks we used to think needed a human mind: writing, recognising images, making predictions.
Large language model (LLM)	The specific kind of AI behind chat tools. A very large prediction engine trained on enormous amounts of text.
Prompt	Everything you type to the model. Your brief, your question, your instructions.
Token	The small chunk of text a model actually reads. Roughly three quarters of a word in English. Prompts and answers are measured in tokens.
Context window	The maximum amount of text the model can consider at once: your messages, its replies, everything. Outside the window, it knows nothing.
Hallucination	When the model states something false with complete confidence. It predicts plausible text; it does not check truth.
Temperature	A dial for predictability. Low means consistent and safe; high means more creative and more random.
System prompt	Standing instructions set above the conversation by whoever runs the tool. Defines the model's role and rules before you type a word.
Chain of thought	Asking the model to reason step by step before concluding. Improves answers on judgement and calculation tasks.
Few shot	Giving the model several examples of the pattern you want, then asking it to continue the pattern.
Zero shot	Asking for something with no examples at all. Fine for simple tasks, unreliable for unusual formats.
Persona	The role you assign the model: "you are a cautious accountant". Changes the style and emphasis it draws on.
Prompt engineering	The craft of writing prompts that reliably produce what you want. Everything in this book.
Fine tuning	Extra training that teaches a model a specific style or domain. Expensive, slow, and usually unnecessary when good prompting will do.
RAG (retrieval augmented generation)	Bolting a search of your own documents onto the model, so it answers from your files rather than only its training.

TERM	WHAT IT ACTUALLY MEANS
Agent	A model given tools and permission to act: search the web, run code, send emails. Powerful, and needs tight limits.
Model	A particular version of the engine, such as a named release from OpenAI, Anthropic or Google. Newer is not automatically better for your task.
Parameters	The internal numbers the model learned in training. More parameters generally means a bigger, more capable model.
Training data	The text the model learned from. Its knowledge, biases and blind spots all come from here.
Inference	The act of generating an answer, as opposed to training. Every reply you get is inference.
Jailbreak	A prompt designed to make a model ignore its safety rules. Curiosity about these is natural; building your work on them is foolish.
Guardrails	The rules and filters that stop a model producing harmful or off limits content.
Multimodal	A model that handles more than text: images, audio, video in or out.
Latency	The delay between sending your prompt and the answer arriving. Longer prompts and bigger models mean more latency.
Hallmark phrases	The verbal tics that reveal unedited AI text: "in today's fast paced world", "delve", "unlock". Ban them in your prompts.

About the Author and Clone Centre

> About Joseph

I am Joseph Conlan, an AI consultant and the founder of Clone Centre, based in Liverpool in the UK. I have spent my career at the point where technology meets real businesses, and I have advised billion dollar VC funds on where AI is genuinely heading, as opposed to where the headlines say it is heading.

My conviction is simple: the winners of the next decade will not be the companies with the most AI. They will be the companies whose people communicate best with it. That skill is learnable, it is cheap, and it starts with the contents of this book.

THE CLONE CENTRE PROMPT GUIDEBOOK · CLONECENTRE.AI

If this guide helped you, that is the whole point. If you want to go further, that is what Clone Centre is for.

> About Clone Centre

Clone Centre is an AI consultancy and engineering studio. We build human first AI: systems that serve people, never the other way round. For our clients that means practical work with real returns: AI strategy, custom built tools, automation that removes drudgery, and the training that makes it all stick, with compliance and governance built in from day one rather than bolted on after.

We work with founders, leadership teams and regulated businesses who want the benefits of AI without the recklessness. If that sounds like you, come and talk to us.

Find us at **clonecentre.ai**.

KEEP THIS BOOK CLOSE

Read Chapters 1 to 4 once, and they will change how you think. Keep Chapter 5 within reach of your keyboard, and it will change what you produce. And when you write a prompt that works beautifully, save it. That is your library beginning, and it starts today.

CLONE CENTRE

Human first AI, built in Liverpool.

clonecentre.ai